Tejas Gokhale

CONTACT
Email: gokhale@umbc.edu Website: https://www.tejasgokhale.com
Mail: 1000 Hilltop Circle, ITE 214, Baltimore MD 21250

APPOINTMENT

Assistant Professor
Department of Computer Science & Electrical Engineering
University of Maryland, Baltimore County

Affiliated Faculty, UMBC AI Center

RESEARCH AREA

Robust computing for perception, communication, learning, and reasoning.
Computer Vision, Machine Learning, Reliable AI, Multimodal Learning

EDUCATION

Doctor of Philosophy, Arizona State University 2023
School of Computing and Augmented Intelligence
Advisors: Yezhou Yang, Chitta Baral
Thesis: Towards Reliable Semantic Vision

Master of Science, Carnegie Mellon University 2017
Department of Electrical and Computer Engineering
Mentor: Aswin Sankaranarayanan

Bachelor of Engineering (Honours), BITS Pilani 2015
Department of Electrical and Electronics Engineering

EMPLOYMENT HISTORY

Microsoft Research Summer 2022
Research Intern, Adaptive Systems and Interaction Group
Mentors: Hamid Palangi, Besa Nushi, Vibhav Vineet, Eric Horvitz

Lawrence Livermore National Laboratory Summer 2021, 2020
Research Scholar, Machine Intelligence Group
Mentors: Rushil Anirudh, Jay Thiagarajan, Bhavya Kailkhura

Arizona State University
Graduate Research Associate, School of Computing and AI 2018–2023
Graduate Teaching Associate, School of Computing and AI 2018–2020

Snap Research Summer 2018
Research Intern, Computational Imaging Group
Mentors: Guru Krishnan, Shree Nayar

Carnegie Mellon University 2017–2018
Graduate Student Researcher, Dept. of Electrical and Computer Engineering

ST Microelectronics Fall 2014
Intern, High Speed Links Group

Steel Authority of India Limited Summer 2013
Summer Intern, Bhilai Steel Plant

PUBLICATIONS See my Google Scholar page for recent updates and citation information.

Conference Proceedings

[C1] Sourajit Saha, Tejas Gokhale. Improving Shift Invariance in Convolutional Neural Networks with Translation Invariant Polyphase Sampling. In Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision. 2025.
https://arxiv.org/abs/2404.07410 WACV 2025

[C2] Maitreya Patel, Abhiram Kusumba, Sheng Cheng, Changhoon Kim, Tejas Gokhale, Chitta Baral, Yezhou Yang. TripletCLIP: Improving Compositional Reasoning of CLIP via Vision-Language Negatives. In Advances in Neural Information Processing Systems. 2024.
(to--appear) NeurIPS 2024

[C3] Agneet Chatterjee, Gabriela Ben Melech Stan, Estelle Guez Aflalo, Sayak Paul, Dhruba Ghosh, Tejas Gokhale, Ludwig Schmidt, Hannaneh Hajishirzi, Vasudev Lal, Chitta Baral, Yezhou Yang. Getting it Right: Improving Spatial Consistency in Text-to-Image Models. In European conference on computer vision. 2024.
https://arxiv.org/abs/2404.01197 ECCV 2024

[C4] Agneet Chatterjee, Yiran Luo, Tejas Gokhale, Chitta Baral, Yezhou Yang. REVISION: Rendering Tools Enable Spatial Fidelity in Vision-Language Models. In European conference on computer vision. 2024.
https://arxiv.org/abs/2408.02231 ECCV 2024

[C5] Agneet Chatterjee, Tejas Gokhale, Chitta Baral, Yezhou Yang. On the Robustness of Language Guidance for Low-Level Vision Tasks: Findings from Depth Estimation. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, pp. 2794-2803. 2024.
https://arxiv.org/abs/2404.08540 CVPR 2024

[C6] Maitreya Patel, Tejas Gokhale, Chitta Baral, Yezhou Yang. ConceptBed: Evaluating Concept Learning Abilities of Text-to-Image Diffusion Models. In Proceedings of the AAAI Conference on Artificial Intelligence, vol. 38, no. 13, pp. 14554-14562. 2024.
https://arxiv.org/abs/2306.04695 AAAI 2024

[C7] Sheng Cheng, Tejas Gokhale, Yezhou Yang. Adversarial Bayesian Augmentation for Single-Source Domain Generalization. In Proceedings of the IEEE/CVF International Conference on Computer Vision, pp. 11400-11410. 2023.
https://arxiv.org/abs/2307.09520 ICCV 2023

[C8] Man Luo, Zhiyuan Fang, Tejas Gokhale, Yezhou Yang, Chitta Baral. End-to-end Knowledge Retrieval for Multi-modal Queries. In Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), pp. 8573-8589. 2023.
https://arxiv.org/abs/2306.00424 ACL 2023

[C9] Tejas Gokhale, Rushil Anirudh, Jayaraman J. Thiagarajan, Bhavya Kailkhura, Chitta Baral, and Yezhou Yang. Improving Diversity with Adversarially Learned Transformations for Domain Generalization. In Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision, pp. 434-443. 2023.
https://arxiv.org/abs/2206.07736 WACV 2023

[C10] Maitreya Patel, Tejas Gokhale, Chitta Baral, and Yezhou Yang. CRIPP-VQA: Counterfactual Reasoning about Implicit Physical Properties via Video Question Answering. In Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, pp. 9856-9870. 2022.
https://arxiv.org/abs/2211.03779 EMNLP 2022

[C11] Tejas Gokhale, Abhishek Chaudhary, Pratyay Banerjee, Chitta Baral, and Yezhou Yang. Semantically Distributed Robust Optimization for Vision-and-Language Inference. In Findings of the Association for Computational Linguistics: ACL 2022, pp. 1493-1513. 2022.
https://arxiv.org/abs/2110.07165 ACL Findings 2022

[C12] Tejas Gokhale, Swaroop Mishra, Man Luo, Bhavdeep Sachdeva, and Chitta Baral. Generalized but not Robust? Comparing the Effects of Data Modification Methods on Out-of-Domain Generalization and Adversarial Robustness. In Findings of the Association for Computational Linguistics: ACL 2022, pp. 2705-2718. 2022.
https://arxiv.org/abs/2203.07653 ACL Findings 2022

[C13] Neeraj Varshney, Pratyay Banerjee, Tejas Gokhale, and Chitta Baral. Unsupervised Natural Language Inference Using PHL Triplet Generation. In Findings of the Association for Computational Linguistics: ACL 2022, pp. 2003-2016. 2022.
https://arxiv.org/abs/2110.08438 ACL Findings 2022

[C14] Yiran Luo, Pratyay Banerjee, Tejas Gokhale, Yezhou Yang, and Chitta Baral. To Find Waldo You Need Contextual Cues: Debiasing Who's Waldo. In 60th Annual Meeting of the Association for Computational Linguistics, ACL 2022, pp. 355-361. 2022.
https://arxiv.org/abs/2203.16682 ACL 2022

[C15] Man Luo, Arindam Mitra, Tejas Gokhale, and Chitta Baral. Improving biomedical information retrieval with neural retrievers. In Proceedings of the AAAI Conference on Artificial Intelligence, vol. 36, no. 10, pp. 11038-11046. 2022.
https://arxiv.org/abs/2201.07745 AAAI 2022

[C16] Pratyay Banerjee, Tejas Gokhale, Yezhou Yang, and Chitta Baral. Weakly supervised relative spatial reasoning for visual question answering. In Proceedings of the IEEE/CVF International Conference on Computer Vision, pp. 1908-1918. 2021.
https://arxiv.org/abs/2109.01934 ICCV 2021

[C17] Pratyay Banerjee, Tejas Gokhale, Yezhou Yang, and Chitta Baral. WeaQA: Weak Supervision via Captions for Visual Question Answering. In Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021, pp. 3420-3435. 2021.
https://arxiv.org/abs/2012.02356 ACL Findings 2021

[C18] Pratyay Banerjee, Tejas Gokhale, and Chitta Baral. Self-Supervised Test-Time Learning for Reading Comprehension. In Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, pp. 1200-1211. 2021.
https://arxiv.org/abs/2103.11263 NAACL 2021

[C19] Tejas Gokhale, Rushil Anirudh, Bhavya Kailkhura, Jayaraman J. Thiagarajan, Chitta Baral, and Yezhou Yang. Attribute-guided adversarial training for robustness to natural perturbations. In Proceedings of the AAAI Conference on Artificial Intelligence, vol. 35, no. 9, pp. 7574-7582. 2021.
https://arxiv.org/abs/2012.01806 AAAI 2021

[C20] Tejas Gokhale, Pratyay Banerjee*, Chitta Baral, and Yezhou Yang. MUTANT: A Training Paradigm for Out-of-Distribution Generalization in Visual Question Answering. In Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP), pp. 878-892. 2020.
https://arxiv.org/abs/2009.08566 EMNLP 2020

[C21] Zhiyuan Fang, Tejas Gokhale, Pratyay Banerjee, Chitta Baral, and Yezhou Yang. Video2 Commonsense: Generating Commonsense Descriptions to Enrich Video Captioning. In Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP), pp. 840-

860. 2020.
https://arxiv.org/abs/2003.05162 EMNLP 2020

[C22] Tejas Gokhale, Pratyay Banerjee*, Chitta Baral, and Yezhou Yang. Vqa-lol: Visual question answering under the lens of logic. In European conference on computer vision, pp. 379-396. Cham: Springer International Publishing, 2020.
https://arxiv.org/abs/2002.08325 ECCV 2020

Journals and Magazines

[J1] Tejas Gokhale. Towards Robust Visual Understanding: A Paradigm Shift in Computer Vision from Recognition to Reasoning. AI Magazine 1–7. 2024.
https://doi.org/10.1002/aaai.12194 AI Magazine

Peer Reviewed Workshop Papers and Extended Abstracts

[W1] Yiran Luo, Joshua Feinglass, Tejas Gokhale, Kuan-Cheng Lee, Chitta Baral, Yezhou Yang. Grounding Stylistic Domain Generalization with Quantitative Domain Shift Measures and Synthetic Scene Images. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops. 2024.
https://arxiv.org/abs/2405.15961 CVPR 2024 Vision Datasets Understanding Workshop

[W2] Tejas Gokhale. Towards Robust Visual Understanding: from Recognition to Reasoning. In Proceedings of the AAAI Conference on Artificial Intelligence, vol. 38, no. 20, pp. 22665-22665. 2024.
https://ojs.aaai.org/index.php/AAAI/article/view/30281 AAAI New Faculty Highlights

[W3] Tejas Gokhale, Joshua Feinglass, and Yezhou Yang. Covariate Shift Detection via Domain Interpolation Sensitivity. In First Workshop on Interpolation Regularizers and Beyond at NeurIPS 2022.
https://openreview.net/pdf?id=YkPjTHZDdm NeurIPS 2022 Interpolation Workshop

[W4] Kuldeep Kulkarni, Tejas Gokhale, Rajhans Singh, Pavan Turaga, Aswin C. Sankaranarayanan. Halluci-Net: Scene Completion by Exploiting Object Co-occurrence Relationships. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops. 2021.
https://arxiv.org/abs/2004.08614 CVPR 2021 AI for Content Creation Workshop

[W5] Tejas Gokhale, Shailaja Sampat, Zhiyuan Fang, Yezhou Yang, and Chitta Baral. Cooking with blocks: A recipe for visual reasoning on image-pairs. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops, pp. 5-8. 2019.
https://arxiv.org/abs/1905.12042 CVPR 2019 Vision Meets Cognition Workshop

[W6] Tejas Gokhale. Vision beyond Pixels: Visual Reasoning via Blocksworld Abstractions. In IJCAI, pp. 6436-6437. 2019.
https://www.ijcai.org/Proceedings/2019/0907.pdf IJCAI 2019 Doctoral Consortium

Technical Reports and Preprints

[P1] Maitreya Patel, Neeraj Varshney, Agneet Chatterjee, Tejas Gokhale, Yezhou Yang, Chitta Baral. Reliability-Checklist: Framework for Comprehensively Evaluating the Reliability of NLP Systems.
https://github.com/Maitreyapatel/reliability-checklist Tech Report

[P2] Tejas Gokhale, Hamid Palangi, Besmira Nushi, Vibhav Vineet, Eric Horvitz, Ece Kamar, Chitta Baral, and Yezhou Yang. Benchmarking Spatial Relationships in Text-to-Image Generation. 2022.
https://arxiv.org/abs/2212.10015 Tech Report

[P3] Ethan Wisdom, Tejas Gokhale, Chaowei Xiao, and Yezhou Yang. Mole Recruitment: Poisoning of Image Classifiers via Selective Batch Sampling.
https://arxiv.org/abs/2303.17080 Tech Report

Books and Monographs

[B1] *Advances in Multimodal Information Retrieval and Generation* Springer
Synthesis Lectures on Computer Vision, ISBN: 978-3-031-57815-1
Man Luo, Tejas Gokhale, Neeraj Varshney, Yezhou Yang, Chitta Baral. [website]

Chapters:

- Man Luo, Tejas Gokhale, Neeraj Varshney, Yezhou Yang, and Chitta Baral. "Transformer-Driven Models for Language, Vision, and Multimodality." In Advances in Multimodal Information Retrieval and Generation, pp. 11-34. Cham: Springer International Publishing, 2024.
- Man Luo, Tejas Gokhale, Neeraj Varshney, Yezhou Yang, and Chitta Baral. "Multimodal Information Retrieval." In Advances in Multimodal Information Retrieval and Generation, pp. 35-91. Cham: Springer International Publishing, 2024.
- Man Luo, Tejas Gokhale, Neeraj Varshney, Yezhou Yang, and Chitta Baral. "Multimodal Content Generation." In Advances in Multimodal Information Retrieval and Generation, pp. 93-134. Cham: Springer International Publishing, 2024.
- Man Luo, Tejas Gokhale, Neeraj Varshney, Yezhou Yang, and Chitta Baral. "Retrieval Augmented Modeling." In Advances in Multimodal Information Retrieval and Generation, pp. 135-157. Cham: Springer International Publishing, 2024.

Ph.D. Dissertation

[T1] Tejas Gokhale. 2023. Towards Reliable Semantic Vision. Order No. 30426752, Arizona State University. https://www.proquest.com/docview/2813822780

Intellectual Property

[IP1] *Automated Evaluation of Spatial Relationships in Images* (US Patent App. 18/198,593)

[IP2] *Systems, Methods, and Apparatuses for Implementing Improved Diversity using Adversarially Learned Transformations for Domain Generalization* (US Patent App. 63/468,653)

Teaching

Instructor, UMBC

CMSC 475/675 Neural Networks [website]	Spring 2025
CMSC 491/691 Robust Machine Learning [website]	Fall 2024
CMSC 491/691 Computer Vision [website]	Spring 2024, Fall 2023
CMSC 898 Pre-Doctoral Candidacy Research	Spring 2025, Fall 2024, Spring 2024
CMSC 799 Master's Thesis Research	Spring 2025, Fall 2024
CMSC 699 Independent Study	Spring 2025, Fall 2024, Spring 2024, Fall 2023

Graduate Teaching Associate, Arizona State University

CSE310: Data Structures & Algorithms	Spring 2020
CSE408: Multimedia Information Systems	Spring 2019
CSE110: Introduction to Programming,	Fall 2018

Guest Lecturer, Arizona State University

CSE598, Perception in Robotics	Spring 2022
CSE408, Multimedia Information Systems	Spring 2019

Student Instructor, BITS Pilani Goa Campus
CTE: Advanced Image Processing — Spring 2015

Funding

UMBC Strategic Awards for Research Transitions (START) — 2024-25
"A Framework for Quantifying Typicality of AI-Generated Images" — $25,000

Maryland Procurement Office (via Johns Hopkins University) — 2024
"Modular Natural Language Understanding" — my share: ∼$43k

UMBC Summer Research Faculty Fellowship (SURFF) — 2024
"Improving the Continual Learning Ability of Visual Recognition Systems via Targeted Unlearning" — $8,000

Microsoft Research Accelerate Foundation Models Academic Research — 2024
Cloud Computing and OpenAI Credits — $20,000

Google Cloud — AY 2023-24
Education Credits — $2350

Students

PhD

- Sourajit Saha — Ph.D. CS [current], UMBC
- Zhiwei Zhang — Ph.D. CS [current], UMBC
- Shivanand Kundargi — Ph.D. CS [current], UMBC
- Jordan Turley — Ph.D. CS [current], UMBC
- Dylan Lang — Ph.D. CS [current], UMBC

MS Thesis

- Neel Patel — M.S. CS [current], UMBC

PhD (as Committee Member)

- Mark Jarzynski (advisor: Marc Olano) — Ph.D. CS [current], UMBC
- Sheng Cheng (advisor: Yezhou Yang) — Ph.D. CS [current], ASU
- Maitreya Patel (advisor: Yezhou Yang & Chitta Baral) — Ph.D. CS [current], ASU
- Agneet Chatterjee (advisor: Yezhou Yang & Chitta Baral) — Ph.D. CS [current], ASU

MS Thesis (as Committee Member)

- Naomi Angela Tack (advisor: Don Engel) — M.S. CS 2024, UMBC

Undergraduate

- UMBC CWIT Scholar: Chloe Wood — 2024-25
- UMBC CWIT Scholar: Danielle Burton — 2023-24

MS (Thesis) Mentees (before UMBC)

- Maitreya Patel (see publication [C10]) — M.S. CS 2022, ASU [thesis]
- Abhishek Chaudhary (see publication [C11]) — M.S. CS 2021, ASU [thesis]

Undergraduate Mentees (before UMBC)

- ASU FURI Program: Mertay Dayanc — B.S CS, 2020
- ASU CS Capstone Project: Paul Butler, Jace Lord, Aashwin Ranjan, Sagarika Pannase, William Tith — 2019-20

Presentations

(Invited Talk), UMBC Information Systems Seminar — 11/2024
"Cognitive Vision: Concepts, Contexts, and Semantics"

(Tutorial), European Conference on Computer Vision — 10/2024
"Evaluation and Benchmarking for Text-to-Image Models"

(Lightning Talk), IARPA Video-LINCS Proposers Day — 02/2024
"Robust Visual Understanding: Knowledge-Guided and Multimodal Reasoning"

(Tutorial), Winter Conference on Applications of Computer Vision — 01/2024
"Challenges with Evaluation of Text-to-Image Models" [website]

(Invited Talk), PRG Seminar, UMIACS (University of Maryland) — 11/2023
"Robust Visual Understanding in the Multimodal Era"

(Invited Talk) *"Towards Reliable Semantic Vision"* — Spring 2023
Rochester Institute of Technology (02/23), Binghamton University (03/23), University of Maryland Baltimore County (03/23), Indiana University (03/23), Case Western Reserve University (03/23), Colorado School of Mines (03/23), Temple University (04/2023)

(Tutorial), Winter Conference on Applications of Computer Vision — 01/2023
"Semantic Data Engineering for Robustness Under Multimodal Settings" [website]

(Invited Talk) University of Illinois at Chicago — 10/2022
"Robust Semantic Vision"

(Invited Talk) Microsoft Research — 10/2022
"Benchmarking Spatial Relationships in Text-to-Image Generation"

(Doctoral Consortium) CVPR, New Orleans — 06/2022
"Discovering Transformations for Generalization in Semantic Vision"

(Guest Lecture) Arizona State University CSE 598 — 03/2022
"Introduction to Generalization in Semantic Vision"

(Invited Talk) Arizona State University ML Club — 09/2021
"Robust Visual Understanding"

(Doctoral Consortium), IJCAI, Macao — 08/2019
"Vision Beyond Pixels"

(Tutorial) Telluride Neuromorphic Cognition Engineering Workshop, — 07/2019
"Reasoning about Objects and Actions via Block-Play"

(Invited) Birla Institute of Technology and Science (BITS Pilani) — 04/2018
"Deep Learning Methods in Imaging and Computer Vision"

Media

Alum inspires next generation of computer vision researchers — 10/2024
ASU Full Circle

Frontiers of multimodal learning: A responsible AI approach — 09/2023
Microsoft Research Blog

CASC research in ML robustness debuts at AAAI conference — 02/2021
News and Press, LLNL Computing

HuggingFace and Intel release a solution for high-fidelity text and image consistency
NetEase (163.com), China — 04/2024

Academic Service

Area Chair:

- Winter Applications of Computer Vision (WACV) 2025
- Advances in Neural Information Processing Systems (NeurIPS) 2024
- Association for Computational Linguistics (ACL) Rolling Review 2024
- North American Chapter of the ACL (NAACL) 2024
- Empirical Methods in Natural Language Processing (EMNLP) 2024

Reviewer / Program Committee:

- Conference on Computer Vision and Pattern Recognition (CVPR) 2023-24
- International Conference on Computer Vision (ICCV) 2023
- European Conference on Computer Vision (ECCV) 2022-24
- Winter Conference on Applications of Computer Vision (WACV) 2021–24
- International Conference on Machine Learning (ICML) 2023–24
- Advances in Neural Information Processing Systems (NeurIPS) 2022–24
- International Conference on Learning Representations (ICLR) 2022–25
- AAAI Conference on Artificial Intelligence (AAAI) 2021–24
- Conference on Language Models (COLM) 2024
- Association for Computational Linguistics (ACL) 2021–24
- Empirical Methods in Natural Language Processing (EMNLP) 2021–23
- North American Chapter of the ACL (NAACL) 2021–23
- International Conference on Robotics and Automation (ICRA) 2019–2023
- International Conference on Intelligent Robots and Systems (IROS) 2022
- IEEE Robotics and Automation Letter (RA-L) 2020-24
- IEEE Transactions of Pattern Analysis and Machine Intelligence (T-PAMI) 2024-25
- ACM Transactions of Computing for Healthcare 2024
- ACM Computing Surveys 2024
- Springer Machine Vision and Applications (MVAP) 2020
- Springer Book Proposals Reviewer 2024

Leadership:

- Team Lead, Summer Camp for Applied Language Exploration (SCALE) 2024, JHU Human Language Technology Center of Excellence [Website] Summer 2024
- Organizer, Tutorial on Responsibly Building Generative Models [Website] ECCV'24
- Mentor, Undergraduate Student Consortium (AAAI-UC) AAAI 2024
- Best Student Abstract Award Committee, AAAI 2024
- Organizer, Tutorial on Reliability of Generative Models in Vision [Website] WACV'24
- Organizer, Workshop on Open-Domain Reasoning under Multi-Modal Settings (ODRUM), [Website] [YouTube] CVPR'23
- Organizer, Workshop on Open-Domain Retrieval under Multi-Modal Settings (ODRUM), [Website] [YouTube] CVPR'22
- Organizer, Tutorial on Semantic Data Engineering under Multimodal Settings (SERUM) [Website] WACV'23
- Organizer, 2021 Frontiers of V&L Seminar Series, [Website], [YouTube] ASU

University Service

University Service (at UMBC):

- Course Development, CMSC 475/675: Neural Networks [Website]
- Course Development, CMSC 491/691: Robust Machine Learning [Website]
- Course Development, CMSC 491/691: Computer Vision [Website]
- PPR Seminar: Advances in Perception, Prediction, and Reasoning [Website]
- Graduate Admissions Committee 2023–present
- Department Publicity Committee 2023–present
- Faculty Mentor, Center for Women in Technology 2023–present
- Undergraduate Student Advisor 2023–present
- CSEE Lightning Talks and Open House Fall 2023, Fall 2024
- Reviewer, CSEE Research Day Spring 2024
- Reviewer, COEIT Cybersecurity Research and Education Proposals Fall 2024
- Member, Asian and Asian American Faculty Staff Council 2023–present

University Service (at ASU):

- Founder, Summer Vision Reading Group, ASU [Website]
- Course Development, CSE591: Frontier Topics in Vision & Language [website] Spring 2021, ASU [YouTube]
- Volunteer, 2019 Southwest Robotics Symposium, Tempe AZ
- Volunteer, International Conference on Machine Learning 2020, Virtual
- Founding Advisor, ASU Machine Learning Club, ASU
- Award Reviewer, GPSA Teaching Award Reviewer ASU
- Mentor, Graduate Student Mentorship Program, ASU
- Project Mentor, CSE598 - Perception in Robotics, ASU Spring 2022
- Project Mentor, CSE576 - Natural Language Processing, ASU Fall 2018

Awards

CVPR 2024 VDU Workshop, Best Paper Award 2024
Research Excellence Award, ASU GPSA 2022
Outstanding Mentor Award, ASU GPSA 2022
NeurIPS Top Reviewer NeurIPS 2022
CVPR 2022 Doctoral Consortium CVPR 2022
ICLR Best Reviewer ICLR 2022
SCAI Doctoral Fellowship (ASU), 2022, 2021, 2020
Engineering Graduate Fellowship, (ASU Engineering) 2023, 2020
ASU GPSA Travel Award for WACV 2023
Graduate College Travel Award (declined) WACV'23, CVPR'22
Graduate College Travel Award (accepted) ICCV'21, EMNLP'20, ECCV'20
IJCAI 2019 Doctoral Consortium IJCAI 2019
Inducted, IEEE Eta Kappa Nu, Sigma Chapter CMU, 2017
National Talent Scholarship, National Council of Educational Research and Training (Govt. of India) 2007–2015

References

Available upon request