Tejas Gokhale

tgokhale@asu.edu tejasgokhale.com Google Scholar

RESEARCH INTERESTS

My mission is to research and develop robust and reliable AI systems by leveraging the complex interactions between vision and language. I work at the wonderful intersection of machine learning, computer vision, and natural language processing. My domain expertise lies in devising adversarial machine learning algorithms, semantic data engineering techniques, and evaluation protocols for out-of-distribution environments.

EDUCATION

Doctor of Philosophy, Arizona State University 2018–present
School of Computing and Augmented Intelligence
Advisors: Yezhou Yang, Chitta Baral

Master of Science, Carnegie Mellon University 2017
Department of Electrical and Computer Engineering
Advisor: Aswin Sankaranarayanan

Bachelor of Engineering (Honours), Birla Institute of Technology and Science 2015
Department of Electrical and Electronics Engineering

RESEARCH EMPLOYMENT

Microsoft Research Summer 2022
Research Intern, Adaptive Systems and Interaction Group
Mentors: Hamid Palangi (+Besa Nushi, Vibhav Vineet, Eric Horvitz)

Lawrence Livermore National Laboratory Summer 2021
Research Scholar, Machine Intelligence Group
Mentor: Rushil Anirudh (+Jay Thiagarajan, Bhavya Kailkhura)

Lawrence Livermore National Laboratory Summer 2020
Research Scholar, Machine Intelligence Group
Mentor: Rushil Anirudh (+Jay Thiagarajan, Bhavya Kailkhura)

Arizona State University 2018–2023
Graduate Research Associate, School of Computing and AI
Yezhou Yang + Chitta Baral

Snapchat Research Summer 2018
Research Intern, Computational Imaging Group
Mentors: Guru Krishnan + Shree Nayar

Carnegie Mellon University 2017–2018
Graduate Student Researcher, Image Science Lab
Advisor: Aswin Sankaranarayanan

ST Microelectronics India Fall 2014
Intern, High Speed Links Group

PUBLICATIONS (REVERSE CHRONOLOGICAL)

My work has been published at AAAI (h5-index: 180); computer vision venues: ICCV (h5-index: 239), ECCV (h5-index: 186), WACV (h5-index: 76); NLP venues: ACL (h5-index: 169), EMNLP (h5-index: 154), NAACL (h5-index: 105).

⊙ Conference Proceedings

[C1] *Improving Diversity with Adversarially Learned Transformations for Domain Generalization*
T. Gokhale, R. Anirudh, J. Thiagarajan, B. Kailkhura, C. Baral, Y. Yang
https://arxiv.org/abs/2206.07736 to appear in WACV 2023

[C2] *CRIPP-VQA: Counterfactual Reasoning about Implicit Physical Properties via Video Question Answering*
M. Patel, **T. Gokhale**, C. Baral, Y. Yang
https://arxiv.org/abs/2211.03779 to appear in EMNLP 2022

[C3] *Semantically Distributed Robust Optimization for Vision-and-Language Inference*
T. Gokhale, A. Chaudhary, P. Banerjee, C. Baral, Y. Yang
https://arxiv.org/abs/2110.07165 ACL Findings 2022

[C4] *Generalized but not Robust? Comparing the Effects of Data Modification Methods on Out-of-Domain Generalization and Adversarial Robustness*
T. Gokhale, S. Mishra, M. Luo, B. Sachdeva, C. Baral
https://arxiv.org/abs/2203.07653 ACL Findings 2022

[C5] *To Find Waldo You Need Contextual Cues: Debiasing Who's Waldo*
Y. Luo, P. Banerjee, **T. Gokhale**, Y. Yang, C. Baral
https://arxiv.org/abs/2203.16682 ACL 2022

[C6] *Unsupervised Natural Language Inference Using PHL Triplet Generation*
N. Varshney, P. Banerjee, **T. Gokhale**, C. Baral
https://arxiv.org/abs/2110.08438 ACL Findings 2022

[C7] *To Find Waldo You Need Contextual Cues: Debiasing Who's Waldo*
Y. Luo, P. Banerjee, **T. Gokhale**, Y. Yang, C. Baral
https://arxiv.org/abs/2203.16682 ACL 2022

[C8] *Improving Biomedical Information Retrieval with Neural Retrievers*
M. Luo, A. Mitra, **T. Gokhale**, C. Baral
https://arxiv.org/abs/2201.07745 AAAI 2022

[C9] *Weakly Supervised Relative Spatial Reasoning for Visual Question Answering*
P. Banerjee, **T. Gokhale**, Y. Yang, C. Baral
https://arxiv.org/abs/2109.01934 ICCV 2021

[C10] *WeaQA: Weak Supervision via Captions for Visual Question Answering*
P. Banerjee, **T. Gokhale**, Y. Yang, C. Baral
https://arxiv.org/abs/2012.02356 ACL Findings 2021

[C11] *Self-Supervised Test-Time Learning for Reading Comprehension*
P. Banerjee, **T. Gokhale**, C. Baral
https://arxiv.org/abs/2103.11263 NAACL 2021

[C12] *Attribute-Guided Adversarial Training for Robustness to Natural Perturbations*
T. Gokhale, R. Anirudh, B. Kailkhura, J. Thiagarajan, C. Baral, Y. Yang
https://arxiv.org/abs/2012.01806 AAAI 2021

[C13] *Mutant: A Training Paradigm for Out-of-Distribution Generalization in Visual Question Answering*
T. Gokhale, P. Banerjee, C. Baral, Y. Yang
https://arxiv.org/abs/2009.08566 EMNLP 2020

[C14] *Video2commonsense: Generating commonsense descriptions to enrich video captioning*
Z. Fang*, **T. Gokhale***, P. Banerjee, C. Baral, Y. Yang
https://arxiv.org/abs/2003.05162 EMNLP 2020

[C15] *VQA-LOL: Visual question answering under the lens of logic*
T. Gokhale, P. Banerjee, C. Baral, Y. Yang
https://arxiv.org/abs/2002.08325 ECCV 2020

⊙ Workshop Proceedings *(2 CVPR, 1 NeurIPS)*

[W1] *Covariate Shift Detection via Domain Interpolation Sensitivity*
T. Gokhale, J. Feinglass, Y. Yang
https://openreview.net/pdf?id=YkPjTHZDdm [SPOTLIGHT] NeurIPS 2022 Interpolate Workshop

[W2] *Halluci-Net: Scene Completion by Exploiting Object Co-occurrence Relationships*
K. Kulkarni, **T. Gokhale**, R. Singh, P. Turaga, A. Sankaranarayanan
https://arxiv.org/abs/2004.08614 AI for Content Creation @ CVPR 2021

[W3] *Cooking With Blocks: A Recipe for Visual Reasoning on Image-Pairs*
T. Gokhale, S. Sampat, Z. Fang, Y. Yang, C. Baral
Long version: https://arxiv.org/abs/1905.12042 Vision Meets Cognition @ CVPR'19

⊙ Preprints

[P1] *Benchmarking Spatial Relationships in Text-to-Image Generation*
T. Gokhale, H. Palangi, B. Nushi, V. Vineet, E. Horvitz, E. Kamar, C. Baral, Y. Yang
https://arxiv.org/abs/2212.10015 in review

[P2] *Poisoning of Image Classifiers via Selective Batch Sampling*
E. Wisdom, **T. Gokhale**, Y. Yang in review

[P3] *End-to-end Knowledge Retrieval for Multi-modal Queries*
M. Luo, Z. Fang, T. Gokhale, Y.Yang, C. Baral in review

⊙ Book Manuscript

[B1] *Advances in Multi-Modal Information Retrieval*
*(*In Preparation) Springer Synthesis Lectures

⊙ Grant Writing

I am actively involved in conceptualizing and writing grant proposals with my advisors. This proposed work builds upon contributions made by my PhD thesis. Note: I am not a PI on these grants.

[1] *Environment-driven Conceptual Learning*
PI: Chitta Baral — *Submitted to DARPA, 2022*

[2] *Decentralized Authorship Attribution*
PI: Chitta Baral — *Submitted to IARPA, 2022*

[3] *An Active Approach for Data Engineering to Improve Vision-Language Tasks*
PI: Yezhou Yang, Co-PI: Chitta Baral — *Funded by NSF, 2021*

INVITED TALKS

Jan'23, *(Tutorial)* "Semantic Data Engineering for Robustness Under Multimodal Settings" — WACV 2023, Hawaii
Oct'22, *(Invited Talk)* "Robust Semantic Vision" — University of Illinois at Chicago
Oct'22, *(Invited Talk)* "Benchmarking Spatial Relationships in Text-to-Image Generation" — Microsoft Research
Mar'22, *(Guest Lecture)* "Introduction to Generalization in Semantic Vision" — ASU CSE 598
Sep'21, *(Invited)* "Robust Visual Understanding", — ASU ML Club
Aug'19, "Vision Beyond Pixels", IJCAI Doctoral Consortium, — IJCAI 2019, Macao
Jul'19, "Reasoning about Objects and Actions via Block-Play", — Telluride 2019
Apr'18, *(Invited)* "Deep Learning Methods in Imaging and Computer Vision", — BITS Goa

TEACHING

Tutorial
SERUM: Semantic Data Engineering for Robustness Under Multimodal Settings — WACV 2023, Hawaii

Teaching Assistant, Arizona State University
CSE310: Data Structures & Algorithms — Spring 2020, ASU
CSE408: Multimedia Information Systems, — Spring 2019, ASU
CSE110: Introduction to Programming, — Fall 2018, ASU

Guest Lecturer
CSE598, Perception in Robotics — Spring 2022, ASU
CSE408, Multimedia Information Systems — Spring 2019, ASU

Course Development
CSE591: Frontier Topics in Vision & Language — [YouTube] [website] Spring 2021, ASU
CTE: Advanced Image Processing, — Spring 2015, BITS Pilani

MENTORING

PhD Students
Ethan Wisdom (see publication [P2]) — Ph.D. CS [current]
Maitreya Patel (see publication [C2]) — Ph.D. CS [current]
Agneet Chatterjee — Ph.D. CS [current]

MS (Thesis) Students
Maitreya Patel (see publication [C2]) — M.S. CS
Abhishek Chaudhary (see publication [C3]) — M.S. CS 2021 [thesis]

Capstone Mentor, mentored five students in projects on visual reasoning — AY 2019-20
UG research mentor (FURI @ ASU), Mertay Dayanc — BS CS, 2020
Project Mentor, CSE598 - Perception in Robotics, ASU — Spring 2022
Project Mentor, CSE576 - Natural Language Processing, ASU — Fall 2018

SERVICE / LEADERSHIP

Program Committee / Conference Reviewer

CVPR: IEEE/CVF Conference on Computer Vision and Pattern Recognition	2023
ICML: International Conference on Machine Learning	2023
NeurIPS: Advances in Neural Information Processing Systems	2022
ICLR: International Conference on Learning Representations	2022
AAAI: AAAI Conference on Artificial Intelligence	2021–2023
ECCV: European Conference on Computer Vision	2022
ACL: Annual Meeting of the Association for Computational Linguistics	2021–2023
EMNLP: Conference on Empirical Methods in Natural Language Processing	2021–2022
NAACL: North American Chapter of the Association for Computational Linguistics	2021–2022
WACV: IEEE Winter Conference on Applications of Computer Vision	2021–2023
ICRA: International Conference on Robotics and Automation	2019–2023
IROS: IEEE/RSJ International Conference on Intelligent Robots and Systems	2022

Journal Reviewer

RA-L: IEEE Robotics and Automation Letter	2020
MVAP: Springer Machine Vision and Applications	2020

Workshop Organizer

2nd ODRUM: Workshop on Open-Domain Reasoning under Multi-Modal Settings, [Website] CVPR'23
1st ODRUM Workshop on Open-Domain Retrieval under Multi-Modal Settings, [Website] [YouTube] CVPR'22

Organizer, 2021 Frontiers of V&L Seminar Series, [Website], [YouTube] ASU
Founder, Summer Vision Reading Group, [Website], multi-university initiative

Volunteer, 2019 Southwest Robotics Symposium, Tempe AZ
Volunteer, International Conference on Machine Learning 2020, Virtual
Advisor, ASU Machine Learning Club, ASU
Award Reviewer, GPSA Teaching Award Reviewer ASU
Mentor, Graduate Student Mentorship Program, ASU
Student Mentor, Peer Mentorship Program BITS Pilani

AWARDS AND RECOGNITION

Research Excellence Award, ASU GPSA	2022
Outstanding Mentor Award, ASU GPSA	2022
NeurIPS Top Reviewer	NeurIPS 2022
CVPR 2022 Doctoral Consortium	CVPR 2022
ICLR Best Reviewer	ICLR 2022
SCAI Doctoral Fellowship (ASU),	2020-2022
Engineering Graduate Fellowship, (ASU Engineering),	2020
ASU GPSA Travel Award	for WACV 2023
Graduate College Travel Award,	WACV'23 (declined), CVPR'22 (declined), ICCV'21, EMNLP'20, ECCV'20
IJCAI 2019 Doctoral Consortium,	IJCAI 2019
Inducted, IEEE Eta Kappa Nu, Sigma Chapter,	CMU, 2017
National Talent Scholarship (Govt. of India),	2007–2015

REFERENCES

Yezhou Yang	Associate Professor	Arizona State University	yz.yang@asu.edu
Chitta Baral	Professor	Arizona State University	chitta@asu.edu
Rushil Anirudh	Research Scientist	Lawrence Livermore National Laboratory	anirudh1@llnl.gov
Eric Horvitz	Chief Scientific Officer	Microsoft	horvitz@microsoft.com
Heni Ben Amor	Associate Professor	Arizona State University	hbenamor@asu.edu